

Integration of Open OnDemand with the Jobstats Job Monitoring Platform

Josko Plazonic¹ and Jonathan Halverson²

¹Research Computing, Princeton University

²Princeton Institute for Computational Science and Engineering, Princeton University

GOOD Conference at Harvard University
March 19, 2025 at 17:10-17:20

Why the Jobstats Job Monitoring Platform?


Multinode CPU fragmentation (96 nodes, 1 CPU-core per node)

JobID	State	Start	Elapsed	NNodes	NCPUS	JobName
1929990	TIMEOUT	2025-03-10T14:19:24	02:05:21	96	96	coffee-+
1930000	TIMEOUT	2025-03-10T14:30:29	05:35:16	5	448	merger-+
1930364	CANCELLED +	2025-03-10T18:28:21	01:08:49	16	1536	magneta+

GPU job without GPU-enabled software (CPU-only PyTorch)

pysocks	pkgs/main/linux-64::pysocks-1.7.1-py312h06a4308_0
python	pkgs/main/linux-64::python-3.12.9-h5148396_0
pytorch	pytorch/linux-64::pytorch-2.5.1-py3.12_cpu_0
pytorch-mutex	pytorch/noarch::pytorch-mutex-1.0-cpu
pyyaml	pkgs/main/linux-64::pyyaml-6.0.2-py312h5eee18b_0

Design of the Jobstats Job Monitoring Platform



```
$ jobstats 46926519
```

```
=====
```

```
Slurm Job Statistics
```

```
=====
```

```
Job ID: 46926519
User/Account: u12345/ee
Job Name: interactive
State: COMPLETED
Nodes: 1
CPU Cores: 4
CPU Memory: 4GB (1GB per CPU-core)
GPUs: 4
QOS/Partition: gpu-short/gpu
Cluster: della
Start Time: Thu Apr 13, 2023 at 10:59 AM
Run Time: 02:05:28
Time Limit: 02:00:00
```

```
Overall Utilization
```

```
=====
```

```
CPU utilization [||||| 8%]
CPU memory usage [||||||||||||||||||||||||||||||||||||| 92%]
GPU utilization [ 0%]
GPU memory usage [| 2%]
```

Detailed Utilization

```
=====
CPU utilization per node (CPU time used/run time)
  della-l03g11: 00:39:04/08:21:52 (efficiency=7.8%)
CPU memory usage per node - used/allocated
  della-l03g11: 3.7GB/4.0GB (941.9MB/1.0GB per core of 4)
GPU utilization per node
  della-l03g11 (GPU 0): 0% <--- GPU was not used
  della-l03g11 (GPU 1): 0% <--- GPU was not used
  della-l03g11 (GPU 2): 0% <--- GPU was not used
  della-l03g11 (GPU 3): 0% <--- GPU was not used
GPU memory usage per node - maximum used/total
  della-l03g11 (GPU 0): 3.2GB/80.0GB (4.0%)
  della-l03g11 (GPU 1): 1.9GB/80.0GB (2.4%)
  della-l03g11 (GPU 2): 694.4MB/80.0GB (0.8%)
  della-l03g11 (GPU 3): 694.4MB/80.0GB (0.8%)
```

Notes

```
=====
* This job did not use 4 of the 4 allocated GPUs. Please resolve this before
  running additional jobs. Wasting resources causes your subsequent jobs to
  have a lower priority. Is the code capable of using multiple GPUs? Please
  consult the documentation for the software. For more info:
    https://researchcomputing.princeton.edu/KB/gpu-computing
```

\$ jobstats 1749336

=====

Slurm Job Statistics

=====

Job ID: 1749336
User/Account: u12345/astro
Job Name: mlogit_bs
State: CANCELLED
Nodes: 1
CPU Cores: 12
CPU Memory: 48GB (4GB per CPU-core)
QOS/Partition: long/all
Cluster: adroit
Start Time: Tue Apr 18, 2023 at 10:29 PM
Run Time: 1-10:46:22
Time Limit: 2-00:00:00

Overall Utilization

=====

CPU utilization	[	8%]
CPU memory usage	[	4%]

Detailed Utilization

=====

CPU utilization per node (CPU time used/run time)
adroit-h11n4: 1-10:40:39/17-09:16:24 (efficiency=8.3%)
CPU memory usage per node - used/allocated
adroit-h11n4: 2.0GB/48.0GB (173.4MB/4.0GB per core of 12)

Notes

- =====
- * **The overall CPU utilization of this job is 8%. This value is low compared to the target range of 90% and above. Please investigate the reason(s) for the low efficiency. For instance, have you conducted a scaling analysis? For more info:**
<https://researchcomputing.princeton.edu/KB/cpu-utilization>
 - * The CPU utilization of this job (8%) is equal to 1 divided by the number of allocated CPU-cores ($1/12=8\%$). This suggests that you may be running a code that can only use 1 CPU-core. If this is true then allocating more than 1 CPU-core is wasteful. Please consult the documentation for the software to see if it is parallelized. For more info:
<https://researchcomputing.princeton.edu/KB/parallel-code>
 - * This job only used 4% of the 48GB of total allocated CPU memory. For future jobs, please allocate less memory by using a Slurm directive such as `--mem-per-cpu=1G` or `--mem=3G`. This will reduce your queue times and make the resources available for your other jobs. For more info:
<https://researchcomputing.princeton.edu/KB/memory>
 - * For additional job metrics including metrics plotted against time:
<https://myadroit.princeton.edu/pun/sys/jobstats>

Jobstats Features

- **Automatically cancel GPU jobs at 0% utilization**
- Identify and/or email users with the most GPU-hours at 0% utilization
- Identify and/or email users that are over-allocating CPU memory
- and 10 other instances of underutilization

Hi Alan (u12345),

The jobs below have been cancelled because they ran for 2 hours at 0% GPU utilization:


JobID	Cluster	Partition	State	GPUs-Allocated	GPU-Util	Hours
60131148	della	llm	CANCELLED	4	0%	2
60131741	della	llm	CANCELLED	4	0%	2

Replying to this automated email will open a support ticket with Research Computing.

github.com/PrincetonUniversity/jobstats

Open OnDemand Basic Grafana Integration

```
---
v2:
  metadata:
    title: Super Cluster
  ....
  custom:
    grafana:
      host: https://ood.server/grafana
      orgId: 1
      dashboard:
        name: ondemand-clusters
        uid: ZALk3fG9a
        panels:
          cpu: 7
          memory: 9
      labels:
        cluster: cluster
        host: node
        jobid: JobID
```



Open OnDemand Completed Jobs

Files ▾ Jobs ▾ Clusters ▾

🕒 Active Jobs

☰ Completed Jobs

🔧 Job Composer

Completed Jobs

Your Jobs ▾

Start Date:

03 / 11 / 2025 📅

End Date:

03 / 11 / 2025 📅

Submit Query

Your Jobs

Enter Username:

Files are available by clicking on JobID number.

Open OnDemand Completed Jobs

Job 1928348 ("2xIC") [click for full stats]


Memory and CPU Efficiency

Node	Memory Used/Allocated	Memory Per Core	CPU Time Used/Allocated	CPU Efficiency	Cores
stellar-m01g2	4.5 GB / 14.6 GB (30.5%)	2.2 GB / 7.3 GB	3-23:48:19 / 4-00:10:06	99.6%	2
Total:	4.5 GB / 14.6 GB (30.5%)		3-23:48:19 / 4-00:10:06	99.6%	2

GPU Efficiency


Node	GPU	GPU Max Memory Used/Total	Average GPU Utilization
stellar-m01g2	#0	11.4 GB / 40 GB (28.5%)	98.0%
stellar-m01g2	#1	11 GB / 40 GB (27.5%)	98.0%
Total:	2 GPU(s)	22.4 GB / 80 GB (28%)	98.0%

Open OnDemand Helper App



Summary

- `jobstats` command shows CPU/GPU utilization and CPU/GPU memory usage
- **Automatically cancel GPU jobs at 0% utilization**
- Identify and/or send email alerts to users for underutilization
- Used by tens of institutions throughout the world
- Jobstats integrates with OOD making it easy to use



github.com/PrincetonUniversity/jobstats